


AI Endpoints

Powerful, secure, and easy-to-integrate generative AI APIs to enhance your applications.



Mistral AI



Qwen



deepseek



NVIDIA



ultralytics

stability.ai

KEY BENEFITS



DATA SECURITY AND CONFIDENTIALITY

Data security and confidentiality for you and your users is a priority on our platform. Your data will never be used to train or improve our AI models; this is one of our many security guarantees.



A PLATFORM DESIGNED FOR DEVELOPERS

Designed for beginner, intermediate and expert developers, AI Endpoints simplifies integrating advanced AI models into your applications. With comprehensive documentation, standard APIs, and code examples, you can add AI capabilities quickly, without the need for prior AI expertise.



YOUR QUICKEST ACCESS TO THE SMARTEST AI MODELS

Get access to a collection of the market's most advanced and renowned AI models. AI Endpoints provides powerful tools for large language models (LLMs), voice processing, document analysis, and image analysis, whatever your specific needs.



REVERSIBLE TECHNOLOGY WITH NO VENDOR LOCK-IN

With AI Endpoints, you have access to clear, deployable models that can run on your infrastructure or integrate with other cloud services. Using this approach helps you retain IT control—it gives you complete freedom and prevents vendor lock-in.

USES

Conversational AI

Enhance your apps by adding conversational AI that interacts naturally with users in real time. These AI-powered chatbots boost customer engagement, automate customer service, and customise user experiences.

Private code assistant

Use AI-powered voice-to-text for transcription in customer service, meetings, and closed captioning. Providing accessible and customisable audio, text-to-speech improves the user experience by meeting individual needs and preferences.

Voice transcription and interactions

Integrate coding help plugins, such as Continue, into common development environments (IDEs). These tools suggest code, detect errors and automate tasks in real time, while protecting data confidentiality.

PRICING

Models

Category	Params	Name	Input Price	Output Price	Unit
Generic LLM	70B	Llama-3.3-70B-Instruct	0.67 €	0.67 €	Per 1M Tokens
	70B	Llama-3.1-70B-Instruct	0.67 €	0.67 €	
	8x7B	Mixtral 8x7B Instruct	0.63 €	0.63 €	
	13B	Mistral-Nemo-Instruct-2407	0.13 €	0.13 €	
	8B	Llama-3.1-8b	0.10 €	0.10 €	
	7B	Mistral 7b	0.10 €	0.10 €	
Reasoning LLM	70B	DeepSeek-R1-Distill-Llama-70B	0.67 €	0.67 €	
Code LLM	32B	Qwen 2.5 Coder 32b	0.87 €	0.87 €	
	7B	Mamba-Codestral-7B-v0.1	0.19 €	0.19 €	
VLM	70B	Qwen2.5 VL 70b	0.91 €	0.91 €	
	7B	llava-v1.6-mistral-7b-hf	0.29 €	0.29 €	
Embedding		bge-multilingual-gemma2	0.01 €	0.01 €	
		BGE M3	0.01 €	0.01 €	
		BGE Base	0.01 €	0.01 €	

Additional models available in preview for free (Speech, Image generation, Object Detection, NLP, ...)

FEATURES & SPECS



Standard APIs

Popular APIs (like OpenAI) for easy integration.

SLA on API : 99.5 % Uptime for each individual model endpoint.



API key authentication

Easily manage and revoke your API keys.

Rate-limit: 400 Requests Per Minute, limited by Project & Model.



Performance

Achieve high inference performance using OVHcloud’s GPU infrastructure.



Security

Underlying platform with ISO 27000, SOC, and healthcare data certifications.



Confidentiality

Your data is neither reused nor kept



More than 40 models

A constantly updated range of popular, open-weight models.



Lifecycle management

Improve reproducibility with transparent model version management.

Guarantee a decommission lead time: 3 months for all models



Sandbox

Interactively test and explore models for free in a simplified environment